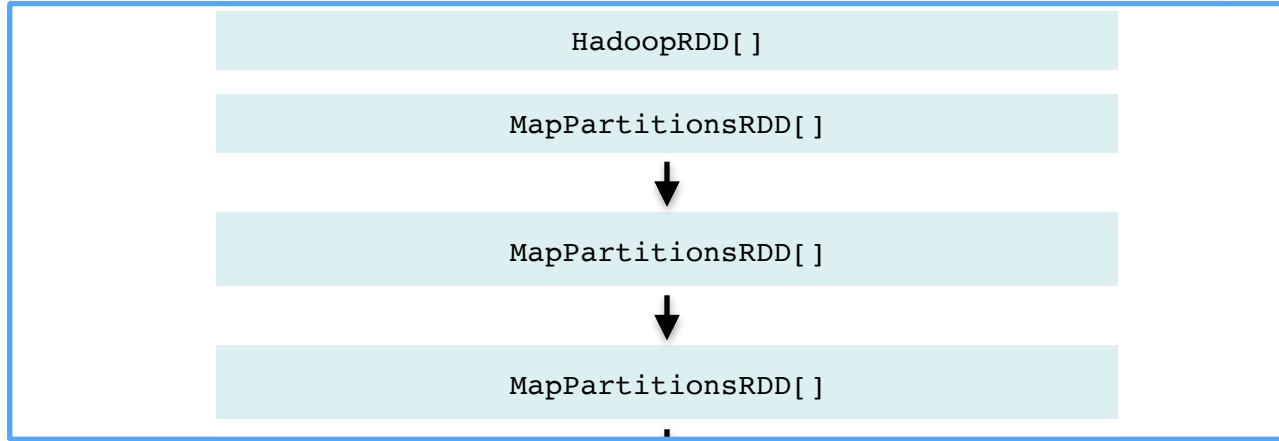


Lesson 4: Spark Internals

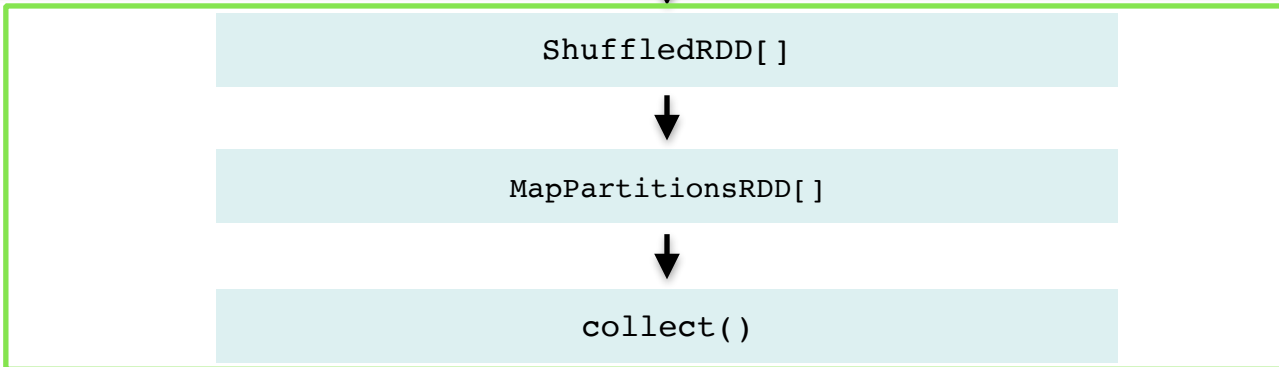
4.6 How Code Runs: Stages, Tasks, and the Shuffle

2. Create Stages

Stage 1



Shuffle forces
Barrier



Stage 2



3. Schedule Tasks

1. Split each **stage** into **tasks** to execute
2. Task is simply a **partition** (data) and **computation** (lambda)
3. Execute all **tasks** in a stage before **continuing**



3. Schedule Tasks

Stage 1

```
sc.textFile('file:///dummy.txt')
```

```
map(line => line.split(" "))
```

```
map(split => (split(0), split(1).toInt))
```

Task 0

Task 1

dummy.txt

```
jon 2  
mary 3  
anna 1  
jon 1
```

```
jesse 1  
mary 5
```

Task 0

partition 0

HadoopRDD

```
map(line => line.split(" "))
```

```
map(split => (split(0), split(1).toInt))
```

Task 1

partition 1

HadoopRDD

```
map(line => line.split(" "))
```

```
map(split => (split(0), split(1).toInt))
```



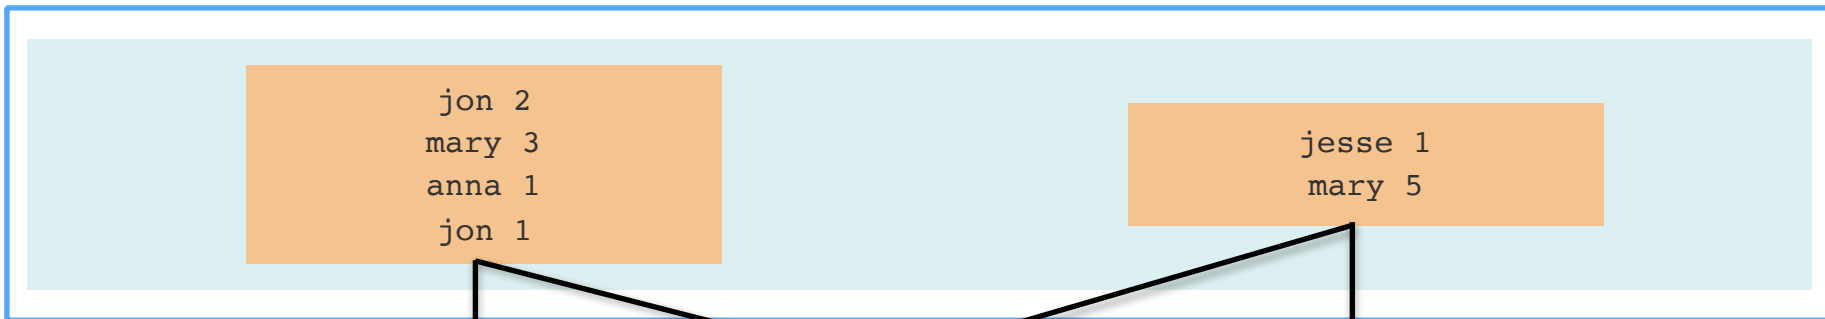
Aside: What's in a Shuffle?

- Redistribute data across **partitions**
- Hashes keys into **buckets**
- **Pull** rather than **push**
- Intermediate files **written to disk** of worker (like Hadoop)

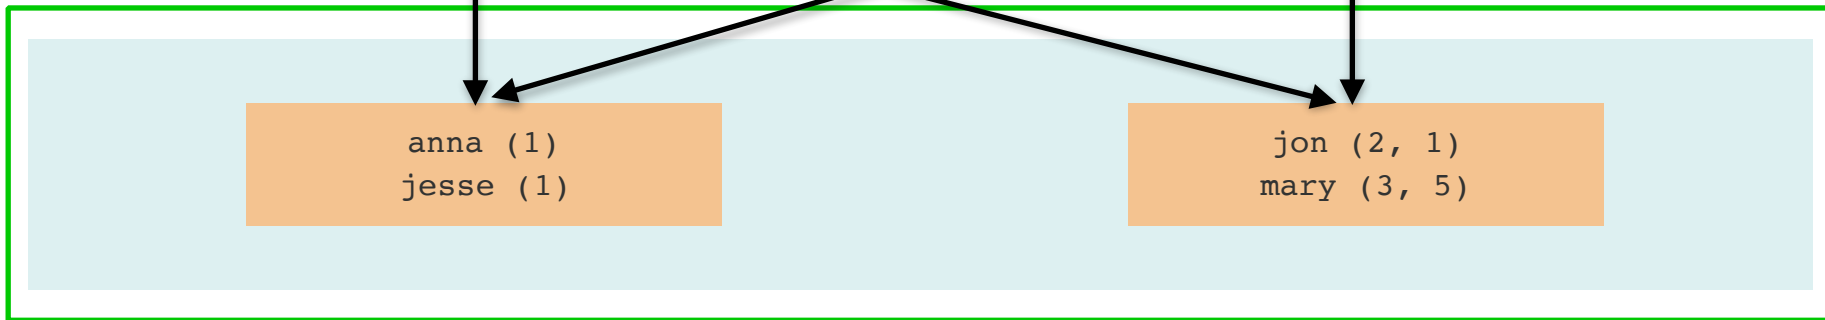


What's in a Shuffle?

Stage 1



groupByKey()



Stage 2



```
file_rdd = sc.textFile('file:///toy_data.txt')
```

```
file_rdd.take(2)
```

```
# => [u'{"Jane": "2"}', u'{"Jane": "1"}']
```

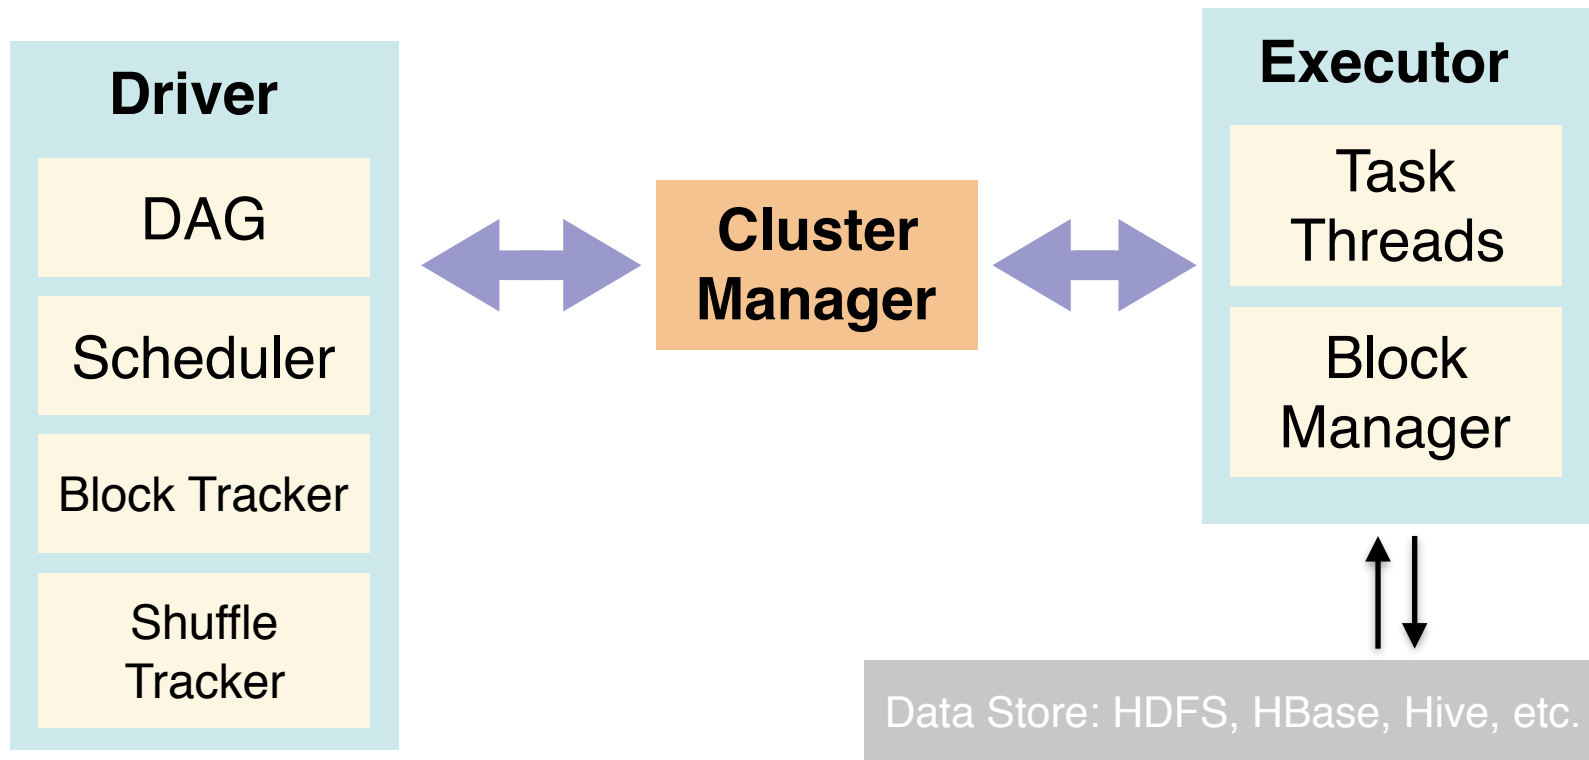
```
tup_rdd = file_rdd.map(lambda line: json.loads(line)) \
                  .map(lambda d: (d.keys()[0], int(d.values()[0])))
```

```
tup_rdd.groupByKey().mapValues(lambda tup: max(tup.data))
```

```
# => [(u'Jane', 2), (u'Pete', 20), (u'Tyler', 3), (u'Yuki', 5), (u'Duncan', 6)]
```



A Day in the Life of a Spark Application



Scheduler

- RDD and partition DAG as **input**
- Tasks (within stages) to execute as **output**

Roles

- Builds stages to execute
- Submit stages to **Cluster Manager**
- **Resubmit** failed stages when output is lost



Review

- The Spark **driver** sends transformations to **cluster manager**
- The Spark **cluster manager** sends **computation** to the appropriate **data**
- Spark intelligently **pipelines** tasks to batch computations to **avoid** sending data over the **network**
- Certain transformations force a **wide dependency**

